

DataEng S24: Data Transformation In-Class Assignment

Submit: Make a copy of this document and use it to record your results. Store a PDF copy of the document in your git repository along with any needed code. Submit the in-class activity submission form by Friday at 10:00 pm.

A. [MUST] Initial Discussion Questions

Discuss the following questions among your working group members at the beginning of the week and place your own response into this space. If desired, also include responses from your group members.

1. In the lecture we mentioned the benefits of Data Transformation, but can you think of any problems that might arise with Data Transformation?

Data Loss: Transformation processes, especially when modifying data types or aggregating data, can result in data loss.

Introduction of Errors: If the transformation logic is flawed, it may introduce errors or inconsistencies in the data.

2. Should data transformation occur before data validation in your data pipeline or after?

Before Data Validation: If we want to transform raw data into a more usable format for validation, we might perform data transformation before data validation. This approach can make the data more consistent and easier to validate.

After Data Validation: Typically, it's more common to perform data validation first, ensuring the data meets certain criteria and quality standards before transforming it. This way, we avoid propagating bad data through your transformation process. After the data is validated and found to be clean and accurate, we can proceed with transformation.

B. [MUST] Small Sample of TriMet data

Here is sample data for one trip of one TriMet bus on one day (February 15, 2023):

[bc_trip259172515_230215.csv](#) It's in .csv format not json format, but otherwise, the data is a typical subset of the data that you are using for your class project.

We recommend that you use google Colab or a Jupyter notebook for this assignment, though any python environment should suffice.

Use the [pandas.read_csv\(\)](#) method to read the data into a DataFrame.

```
#Mount to drive
import pandas as pd
from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
data = pd.read_csv("/content/drive/MyDrive/bc_trip259172515_230215.csv")
data
```

	EVENT_NO_TRIP	EVENT_NO_STOP	OPD_DATE	VEHICLE_ID	METERS	ACT_TIME	GPS_LONGITUDE	GPS_LATITUDE	GPS_SATELLITES	GPS_HDOP
0	259172515	259172517	15FEB2023:00:00:00	4223	40	20469	-122.648137	45.493082	12	0.7
1	259172515	259172517	15FEB2023:00:00:00	4223	48	20474	-122.648240	45.493070	12	0.8
2	259172515	259172517	15FEB2023:00:00:00	4223	57	20479	-122.648352	45.493123	12	0.8
3	259172515	259172517	15FEB2023:00:00:00	4223	73	20484	-122.648385	45.493262	12	0.7
4	259172515	259172517	15FEB2023:00:00:00	4223	112	20489	-122.648347	45.493582	12	0.8
...
156	259172515	259172530	15FEB2023:00:00:00	4223	5834	21389	-122.677057	45.528040	11	0.8
157	259172515	259172531	15FEB2023:00:00:00	4223	5838	21424	-122.677003	45.528037	6	2.0
158	259172515	259172531	15FEB2023:00:00:00	4223	5858	21429	-122.676765	45.528043	7	1.1
159	259172515	259172531	15FEB2023:00:00:00	4223	5889	21434	-122.676370	45.528047	10	1.1
160	259172515	259172531	15FEB2023:00:00:00	4223	5918	21439	-122.675990	45.528065	11	0.8

161 rows × 10 columns

Next steps: [View recommended plots](#)

C. [MUST] Filtering

Some of the columns in our TriMet data are not generally useful for our class project. For example, our contact at TriMet told us that the EVENT_NO_STOP column is not used and can be safely eliminated for any type of analysis of the data.

Use [pandas.DataFrame.drop\(\)](#) to filter the EVENT_NO_STOP column.

For this in-class assignment we won't need the GPS_SATELLITES or GPS_HDOP columns, so drop them as well.

```
✓ 0s # Use DataFrame.drop() to remove the EVENT_NO_STOP, GPS_SATELLITES, and GPS_HDOP columns
columns_to_drop = ['EVENT_NO_STOP', 'GPS_SATELLITES', 'GPS_HDOP']
data_filtered = data.drop(columns=columns_to_drop)

✓ 0s [13] # Display the first few rows of the filtered DataFrame
print("\nFirst few rows of the DataFrame after dropping EVENT_NO_STOP, GPS_SATELLITES, and GPS_HDOP columns:")
print(data_filtered.head())
```

First few rows of the DataFrame after dropping EVENT_NO_STOP, GPS_SATELLITES, and GPS_HDOP columns:

	EVENT_NO_TRIP	OPD_DATE	VEHICLE_ID	METERS	ACT_TIME	\
0	259172515	15FEB2023:00:00:00	4223	40	20469	
1	259172515	15FEB2023:00:00:00	4223	48	20474	
2	259172515	15FEB2023:00:00:00	4223	57	20479	
3	259172515	15FEB2023:00:00:00	4223	73	20484	
4	259172515	15FEB2023:00:00:00	4223	112	20489	

	GPS_LONGITUDE	GPS_LATITUDE
0	-122.648137	45.493082
1	-122.648240	45.493070
2	-122.648352	45.493123
3	-122.648385	45.493262
4	-122.648347	45.493582

Next, start over and this time try filtering these same columns using the `usecols` parameter of the `read_csv()` method.

```
✓ 0s [22] data_path = "/content/drive/MyDrive/bc_trip259172515_230215.csv"

✓ 1s [23] # Specify the columns to include in the DataFrame
columns_to_include = ['EVENT_NO_TRIP', 'OPD_DATE', 'VEHICLE_ID', 'METERS', 'ACT_TIME', 'GPS_LONGITUDE', 'GPS_LATITUDE']

✓ 0s [24] # Read the CSV file and include only the specified columns
df = pd.read_csv(data_path, usecols=columns_to_include)

✓ 0s [25] # Display the first few rows of the filtered DataFrame
print("First few rows of the filtered DataFrame:")
print(df.head())
```

First few rows of the filtered DataFrame:

	EVENT_NO_TRIP	OPD_DATE	VEHICLE_ID	METERS	ACT_TIME	\
0	259172515	15FEB2023:00:00:00	4223	40	20469	
1	259172515	15FEB2023:00:00:00	4223	48	20474	
2	259172515	15FEB2023:00:00:00	4223	57	20479	
3	259172515	15FEB2023:00:00:00	4223	73	20484	
4	259172515	15FEB2023:00:00:00	4223	112	20489	

	GPS_LONGITUDE	GPS_LATITUDE
0	-122.648137	45.493082
1	-122.648240	45.493070
2	-122.648352	45.493123
3	-122.648385	45.493262
4	-122.648347	45.493582

Why might we want to filter columns this way instead of using `drop()`?

Using the `usecols` parameter in `pandas.read_csv()` to filter columns directly during the file reading process can offer several advantages over using `DataFrame.drop()` after reading the file:

Performance: When you specify columns to include using `usecols` during file reading, only the specified columns are read from the file and loaded into memory. This can significantly reduce memory usage and improve performance, especially when dealing with large datasets.

Efficiency: Since only the specified columns are read into memory, the dataset is loaded more efficiently, and you avoid the overhead of loading unnecessary columns and then dropping them later.

Faster Data Loading: By only loading the columns you need, the data reading process can be faster because there is less data to process.

Clarity and Maintainability: Specifying the columns to include in the data reading stage makes the code more explicit and easier to understand. It indicates which columns are relevant to the analysis, and the DataFrame is immediately ready for use with only the necessary columns.

Reduced I/O: When you specify the columns to include, the file reading process only retrieves the data for the specified columns, reducing the amount of input/output operations.

Minimize Data Transformation: When you drop columns using `DataFrame.drop()`, you are creating a new DataFrame with the unnecessary columns removed. This transformation may not be necessary if you can directly load only the required columns.

D. [MUST] Decoding

Notice that the timestamp for each breadcrumb record is encoded in an odd way that might make analysis difficult. The breadcrumb timestamps are represented by two columns, `OPD_DATE` and `ACT_TIME`. `OPD_DATE` merely represents the date on which the bus ran, and it should be constant, unchanging for all breadcrumb records for a single day. The `ACT_TIME` field indicates an offset, specifically the number of seconds elapsed since midnight on that day.

We're not sure why TriMet represents the breadcrumb timestamps this way. We do know that this encoding of the timestamps makes automated analysis difficult. So your job is to decode TriMet's representation and create a new "TIMESTAMP" column containing a [`pandas.Timestamp`](#) value for each breadcrumb.

Suggestions:

- Use `DataFrame.apply()` to apply a function to all rows of your DataFrame
- The applied function should input the two to-be-decoded columns, then it should:
 - create a datetime value from the `OPD_DATE` input using `datetime.strptime()`
 - create a timedelta value from the `ACT_TIME`
 - add the timedelta value to the datetime value to produce the resulting timestamp

```

✓ [27] from datetime import datetime, timedelta
0s
# Function to decode OPD_DATE and ACT_TIME and create a pandas.Timestamp
def create_timestamp(row):
    # Convert OPD_DATE to a datetime object
    date = datetime.strptime(row['OPD_DATE'], '%d%b%Y:%H:%M:%S')

    # Convert ACT_TIME to a timedelta object (from seconds since midnight)
    time_offset = timedelta(seconds=row['ACT_TIME'])

    # Combine date and time_offset to create a timestamp
    timestamp = date + time_offset

    return timestamp

✓ [28] # Apply the function to each row and create the TIMESTAMP column
0s
df['TIMESTAMP'] = df.apply(create_timestamp, axis=1)

✓ [29] # Display the first few rows of the DataFrame with the new TIMESTAMP column
0s
print("First few rows of the DataFrame with the new TIMESTAMP column:")
print(df.head())

```

First few rows of the DataFrame with the new TIMESTAMP column:

	EVENT_NO_TRIP	OPD_DATE	VEHICLE_ID	METERS	ACT_TIME \
0	259172515	15FEB2023:00:00:00	4223	40	20469
1	259172515	15FEB2023:00:00:00	4223	48	20474
2	259172515	15FEB2023:00:00:00	4223	57	20479
3	259172515	15FEB2023:00:00:00	4223	73	20484
4	259172515	15FEB2023:00:00:00	4223	112	20489

	GPS_LONGITUDE	GPS_LATITUDE	TIMESTAMP
0	-122.648137	45.493082	2023-02-15 05:41:09
1	-122.648240	45.493070	2023-02-15 05:41:14
2	-122.648352	45.493123	2023-02-15 05:41:19
3	-122.648385	45.493262	2023-02-15 05:41:24
4	-122.648347	45.493582	2023-02-15 05:41:29

E. [MUST] More Filtering

Now that you have decoded the timestamp you no longer need the OPD_DATE and ACT_TIME columns. Delete them from the DataFrame.

```
[31] # Remove the OPD_DATE and ACT_TIME columns
df = df.drop(columns=['OPD_DATE', 'ACT_TIME'])

[32] # Display the first few rows of the DataFrame with the new TIMESTAMP column
print("First few rows of the DataFrame with the new TIMESTAMP column and without OPD_DATE and ACT_TIME columns:")
print(df.head())
```

First few rows of the DataFrame with the new TIMESTAMP column and without OPD_DATE and ACT_TIME columns:						
	EVENT_NO_TRIP	VEHICLE_ID	METERS	GPS_LONGITUDE	GPS_LATITUDE	\
0	259172515	4223	40	-122.648137	45.493082	
1	259172515	4223	48	-122.648240	45.493070	
2	259172515	4223	57	-122.648352	45.493123	
3	259172515	4223	73	-122.648385	45.493262	
4	259172515	4223	112	-122.648347	45.493582	

	TIMESTAMP
0	2023-02-15 05:41:09
1	2023-02-15 05:41:14
2	2023-02-15 05:41:19
3	2023-02-15 05:41:24
4	2023-02-15 05:41:29

F. [MUST] Enhance

Create a new column, called SPEED, that is a calculation of meters traveled per second. Calculate SPEED for each breadcrumb using the breadcrumb's METERS and TIMESTAMP values along with the METERS and TIMESTAMP values for the immediately preceding breadcrumb record.

Utilize the [pandas.DataFrame.diff\(\)](#) method for this calculation. `diff()` allows you to calculate the difference between a cell value and the preceding row's value for that same column. Use `diff()` to create a new `dMETERS` column and then again to create a new `dTIMESTAMP` column. Then use `apply()` (with a lambda function) to calculate $SPEED = dMETERS / dTIMESTAMP$. Finally, drop the unneeded `dMETERS` And `dTIMESTAMP` columns.

Question: What is the minimum, maximum and average speed for this bus on this trip? (Suggestion: use the `Dataframe.describe()` method to find these statistics)


```
✓ [35] # Calculate the difference between rows for METERS and TIMESTAMP
0s df['dMETERS'] = df['METERS'].diff()
df['dTIMESTAMP'] = df['TIMESTAMP'].diff().dt.total_seconds() # Convert timedelta to seconds
```

```
✓ [36] # Calculate SPEED as dMETERS / dTIMESTAMP
0s df['SPEED'] = df['dMETERS'] / df['dTIMESTAMP']
```

```
✓ [37] # Drop the unneeded dMETERS and dTIMESTAMP columns
0s df.drop(columns=['dMETERS', 'dTIMESTAMP'], inplace=True)
```

```
✓ [38] # Display the first few rows of the DataFrame with the SPEED column
0s print("First few rows of the DataFrame with the SPEED column:")
print(df.head())
```

```
First few rows of the DataFrame with the SPEED column:
  EVENT_NO_TRIP  VEHICLE_ID  METERS  GPS_LONGITUDE  GPS_LATITUDE  \
0      259172515         4223      40    -122.648137    45.493082
1      259172515         4223      48    -122.648240    45.493070
2      259172515         4223      57    -122.648352    45.493123
3      259172515         4223      73    -122.648385    45.493262
4      259172515         4223     112    -122.648347    45.493582

  TIMESTAMP  SPEED
0 2023-02-15 05:41:09    NaN
1 2023-02-15 05:41:14     1.6
2 2023-02-15 05:41:19     1.8
3 2023-02-15 05:41:24     3.2
4 2023-02-15 05:41:29     7.8
```

```
✓ [39] df['SPEED'] = df['SPEED'].fillna(method='bfill')
```

```
✓ [41] # Calculate and display the summary statistics for the SPEED column
0s summary_statistics = df['SPEED'].describe()
# Displaying the statistics
print("Summary statistics for SPEED:")
print(summary_statistics)
# Extracting the minimum, maximum, and average speed
min_speed = summary_statistics['min']
max_speed = summary_statistics['max']
average_speed = summary_statistics['mean']
# Display the minimum, maximum, and average speed
print("\nMinimum speed: {:.2f} m/s".format(min_speed))
print("Maximum speed: {:.2f} m/s".format(max_speed))
print("Average speed: {:.2f} m/s".format(average_speed))
```

```
Summary statistics for SPEED:
count    161.000000
mean      7.192254
std       4.429027
min       0.000000
25%       3.800000
50%       6.400000
75%      10.800000
max      17.400000
Name: SPEED, dtype: float64
```

```
Minimum speed: 0.00 m/s
Maximum speed: 17.40 m/s
Average speed: 7.19 m/s
```

G. [SHOULD] Larger Data Set

Here is breadcrumb data for the same bus TriMet for the entire day (February 15, 2023):

[bc_veh4223_230215.csv](#)

Do the same transformations (parts C through F) for this larger data set. Be careful, you might need to treat each trip separately. For example, you might need to find all of the unique values for the EVENT_NO_TRIP column and then do the transformations separately on each trip.

Questions:

What was the maximum speed for vehicle #4223 on February 15, 2023?

Where and when did this maximum speed occur?

What was the median speed for this vehicle on this day?

```
[45] # Drop the unneeded columns
data.drop(columns=['OPD_DATE', 'ACT_TIME', 'GPS_SATELLITES', 'GPS_HDOP'], inplace=True)

[46] # Calculate dMETERS and dTIMESTAMP
data['dMETERS'] = data['METERS'].diff()
data['dTIMESTAMP'] = data['TIMESTAMP'].diff().dt.total_seconds() # Convert timedelta to seconds

[47] # Calculate SPEED as dMETERS / dTIMESTAMP
data['SPEED'] = data['dMETERS'] / data['dTIMESTAMP']

[48] # Drop the unneeded dMETERS and dTIMESTAMP columns
data.drop(columns=['dMETERS', 'dTIMESTAMP'], inplace=True)

[49] # Calculate max speed, and median speed for each trip
max_speed = data['SPEED'].max()
median_speed = data['SPEED'].median()

[50] # Find the row with maximum speed
max_speed_row = data[data['SPEED'] == max_speed].iloc[0]

[51] # Print results
print(f"Maximum speed for vehicle #4223 on February 15, 2023: {max_speed:.2f} m/s")
print(f"Where and when did this maximum speed occur:")
print(f"  GPS Longitude: {max_speed_row['GPS_LONGITUDE']}")
print(f"  GPS Latitude: {max_speed_row['GPS_LATITUDE']}")
print(f"  Timestamp: {max_speed_row['TIMESTAMP']}")

print(f"\nMedian speed for this vehicle on this day: {median_speed:.2f} m/s")

Maximum speed for vehicle #4223 on February 15, 2023: 17.40 m/s
Where and when did this maximum speed occur:
  GPS Longitude: -122.660822
  GPS Latitude: 45.505452
  Timestamp: 2023-02-15 05:44:49

Median speed for this vehicle on this day: 7.20 m/s
```


H. [ASPIRE] Full Data Set

Here is breadcrumb data for all TriMet vehicles for the entire day (February 15, 2023):

[bc_230215.csv](#)

Do the same transformations (parts C through F) for the entire data set. Again, beware that simple transformations developed in parts C through F probably will need to be modified for the full data set which contains interleaved breadcrumbs from many vehicles.

Questions:

What was the maximum speed for any vehicle on February 15, 2023?

Where and when did this maximum speed occur?

Which vehicle had the fastest mean speed for any single trip on this day? Which vehicle and which trip achieved this fastest average speed?